

EDUCATION:

- **Master of Science in Artificial Intelligence — Northeastern University, Boston, MA | Jan 2025 – Apr 2027**
Relevant Coursework: Machine Learning, Natural Language Processing, Algorithms, Foundation of Artificial Intelligence
 - **Bachelor of Technology in ECE, Sensors and Wearable Technology — VIT, India | Jul 2018 – Jun 2022**
Relevant Coursework: Artificial Intelligence, Neural Networks and Deep Learning, Data Analysis, Data Structures and Algorithms, Calculus, Applied Linear Algebra, Essentials of Machine Learning.
-

SKILLS:

Languages: Python, Java, R, SQL, Shell Scripting

AI/ML: Deep Learning, Reinforcement Learning, NLP, Transformer Architectures, RAG, Conversational AI, Multi-Agent Systems, Model Evaluation, Prompt Engineering.

Frameworks & Libraries: FastAPI, Flask, PyTorch, TensorFlow, Keras, LangChain, LangGraph, Hugging Face Transformers, React, Next.js, OpenCV, Postman, Pygame

Data & Cloud: MySQL, Snowflake, AWS, Tableau, Power BI, Informatica IICS, Hadoop, GCP, ChromaDB, FAISS, Git, CI/CD,

Tools & Concepts: Apache Spark, Databricks, Docker, JIRA, Confluence, Simli, WebRTC, Vector Search, Workflow Automation, Semantic Search, Statistical Analysis, Cursor.

WORK EXPERIENCE:

Generative AI Product Development Fellow | Burnes Center for Social Change, Northeastern University Jan 2026 – Present

- Designed and developed a privacy-first AI feedback platform for InnovateUS that enables government employees to share post-training reflections through voice and text interviews for public-sector impact evaluation.
- Architected a three-stage LLM pipeline for vagueness classification, adaptive follow-up generation, and structured data extraction aligned with the Kirkpatrick evaluation model.
- Built and integrated a React/Next.js frontend with a Python backend using OpenAI Realtime API and WebRTC for conversational voice interaction, session tracking, and structured export workflows.

Python Developer – Level 2 | IICS Developer | Wipro Technologies, Bengaluru, India

Apr 2022 – Oct 2024

- Developed and deployed an internal **Flask-based LSTM-powered error prediction system** trained on historical logs and enriched with NLP-driven rule-based matching to predict and classify recurring issues.
 - Built a knowledge base for error resolution and integrated REST APIs with SQL databases for unified incident tracking, reporting, and multi-environment deployment, reducing manual investigation time by 40%.
 - Automated recurring workflows, file operations, and housekeeping tasks using Python and Shell scripting, saving approximately 20 developer hours per week while improving reliability and platform support operations.
 - Designed and maintained parameterized ETL workflows and reusable mappings in Informatica IICS, supported CI/CD-based job migration, and built Tableau dashboards for platform monitoring, ETL performance, and proactive decision-making.
-

PROJECTS:

ResuMate AI – Multi-Agent AI Hiring Platform | [GitHub](#) | [Demo](#) | Jan 2026 – Mar 2026

(Python, FastAPI, GPT-4o, LangChain, ChromaDB, LiveKit)

- Built a full-stack multi-agent hiring platform with five specialized AI agents for resume analysis, candidate evaluation, research, technical assessment, and career advising, coordinated through a custom Plan → Execute → Reflect → Output orchestration framework.
- Built a real-time AI video interview system using LiveKit WebRTC, OpenAI Realtime API, and Simli avatar rendering, enabling natural voice-to-voice interviews with automated proctoring features such as face detection, eye tracking, and tab monitoring.
- Implemented RAG-based resume analysis and contextual candidate chat using LangChain, ChromaDB, and OpenAI embeddings, along with a candidate-facing AI advisor supporting interview preparation, resume coaching, and career guidance.

Sentence-Level Transformer for Next-Sentence Prediction | [GitHub](#) | Aug 2025 – Dec 2025

(Python, Transformer, Latent Modeling, NLP)

- Designed a sentence-level language model that performs attention over complete sentences instead of individual tokens, shifting the prediction objective from next-token prediction to next-sentence prediction.
- Implemented a two-stage architecture combining a Sentence-VAE for latent sentence representation with a latent-conditioned Transformer decoder to model discourse-level coherence.
- Evaluated the approach on WikiText-2, analyzing training stability, latent quality, and generative coherence relative to token-level baselines.

Royal Rumble Wrestling Simulation | [GitHub](#) | Feb 2025 – Apr 2025

(Python, PPO, Simulated Annealing, Genetic Algorithms, Pygame)

- Built a **multi-agent reinforcement learning environment** simulating adaptive combat behaviors using **PPO**, fitness-based initiator selection, and dynamic policy updates. Combined **Genetic Algorithms** for evolving wrestler traits with **Simulated Annealing** for probabilistic exploration, visualized through a **turn-based Pygame interface**.